

Repeatability Test

Young-Kyoon Suh

May 31, 2017

1 Description

This document presents two sets of histograms based on the two runs of a program under test, called *INC*, with a set of increasing task lengths from 1 second to 4096 seconds. We would like to see if both of the sets have the same shape for the same task length. If so then, we can say that run repeatability is satisfied in our experiment setting. In this study the used protocol eliminates samples involving infrequent, long-running daemons, which is the first step of EMPv5.

In the protocol, we use taskstats C struct to get measures of a captured process. The taskstat's data is delivered via a netlink socket from the kernel space. The receive buffer for the socket is not robust for many observed processes [1]. Fortunately, there is an average of 95 processes per iteration of a run, which turns out to be fine with the struct. For a much more number of processes, the use of `/proc/[pid]/stat` is preferred, as (i) there are equivalent measures available in the `/proc` filesystem, and (ii) there's little constraint on the use as opposed to taskstats.

Now we show histograms of elapsed time (ET) and process time (PT) of INC via the EMPv5 protocol.

2 Histograms on the First Run

This section exhibits histograms on the first run of INC with its task length increasing from 1 second to 4096 seconds, via SEDONA. The detailed description of the base data is from Table 1.

3 Experiment Notes

Table 1 provides a short description of our experimental runs, on which the following histograms are based.

Machine	Task Length (sec)	Description	Experiment Period	Relevant Histograms
sodb9	INC1~INC64	1000 samples, each	2017-03-02 ~ 2017-03-04	Figs. 1, 2, 5, and 6
sodb9	INC128~INC1024	300 samples, each	2017-03-04 ~ 2017-03-11	Figs. 3 and 7
sodb10	INC2048	300 samples	2017-03-02 ~ 2017-03-09	Figs. 4(a) and 8(a)
sodb12	INC4096	300 samples	2017-02-13 ~ 2017-02-27	Figs. 4(b) and 8(b)

Table 1: Notes on experiment runs used for histograms

3.1 ET

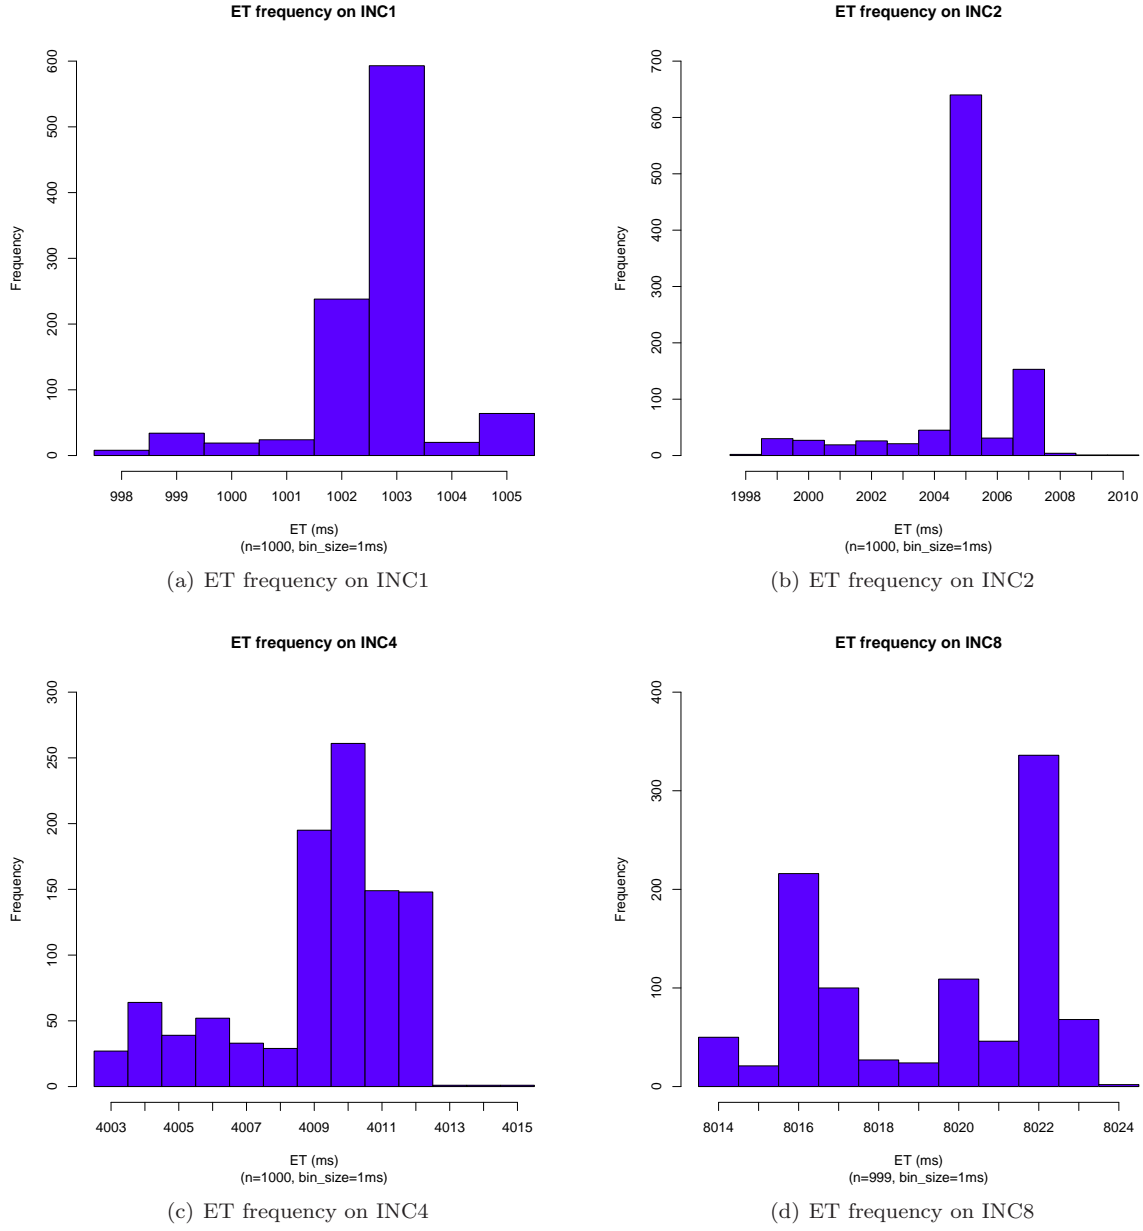





Figure 1: ET Histograms of INC1 ... INC8



(a) ET frequency on INC16

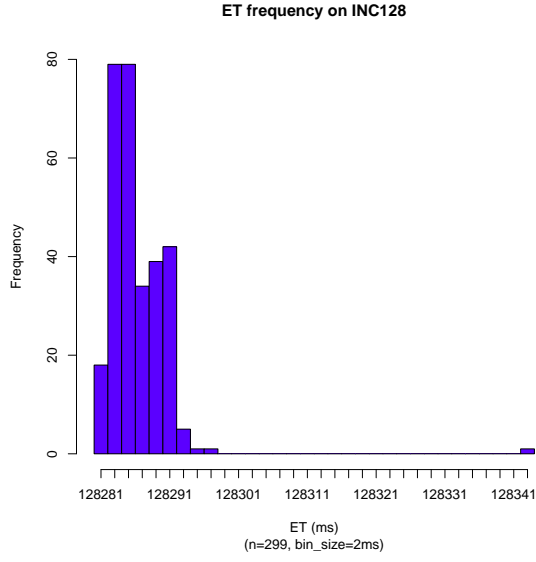


(b) ET frequency on INC32




(c) ET frequency on INC64

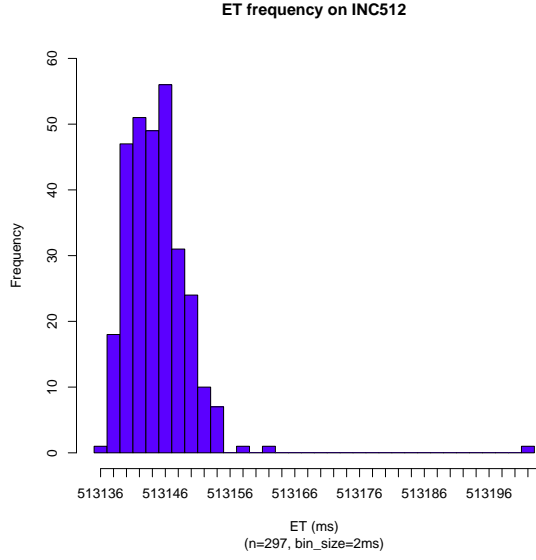
Figure 2: ET Histograms of INC16 ... INC64




(a) ET frequency on INC128



(b) ET frequency on INC256

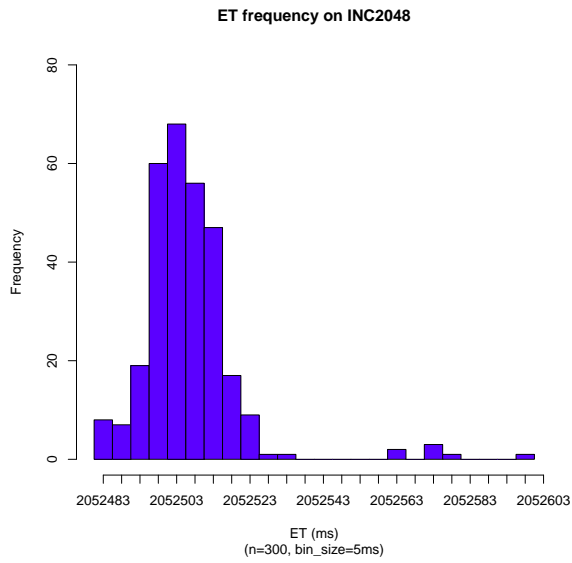


(c) ET frequency on INC512




(d) ET frequency on INC1024

Figure 3: ET Histograms of INC128 ... INC1024



(a) ET frequency on INC2048



(b) ET frequency on INC4096

Figure 4: ET Histograms of INC2048 and INC4096

3.2 PT

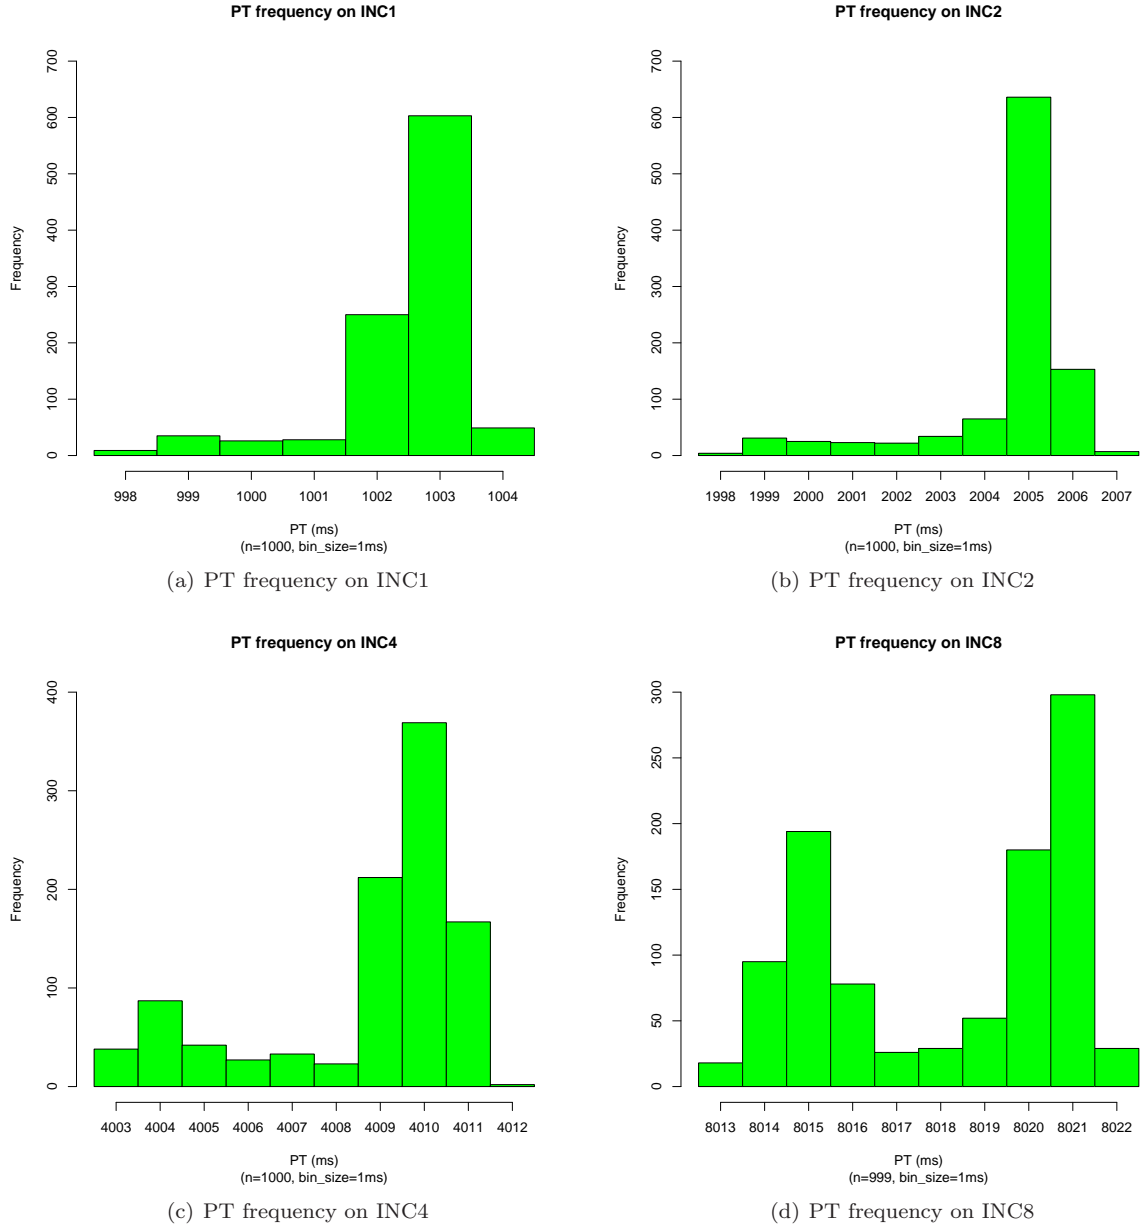





Figure 5: PT Histograms of INC1 ... INC8



(a) PT frequency on INC16




(b) PT frequency on INC32




(c) PT frequency on INC64

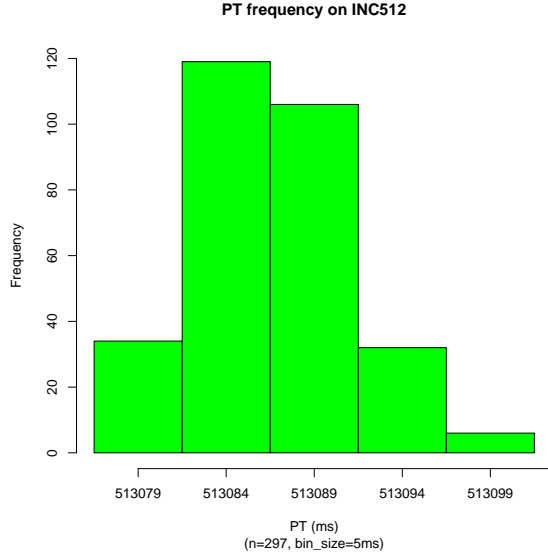
Figure 6: PT Histograms of INC16 ... INC64




(a) PT frequency on INC128



(b) PT frequency on INC256

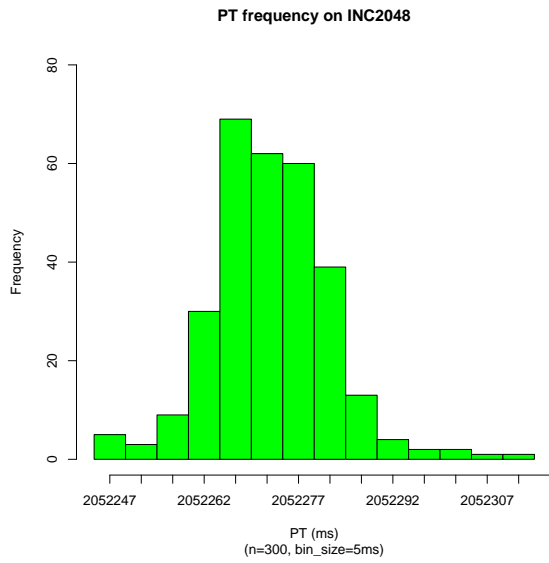


(c) PT frequency on INC512




(d) PT frequency on INC1024

Figure 7: PT Histograms of INC256 ... INC1024



(a) PT frequency on INC2048



(b) PT frequency on INC4096

Figure 8: PT Histograms of INC2048 and INC4096

4 Histograms on the Second Run

This section exhibits histograms on the second run of INC with its task length increasing from 1 second to 4096 seconds, via SEDONA. The detailed description of the base data is from Table 2.

Machine	Task Length (sec)	Description	Experiment Period	Relevant Histograms
sodb9	INC1~INC64	1000 samples, each	2017-03-13 ~ 2017-03-14	Figs. 9, 10, 13, and 14
sodb9	INC128~INC1024	300 samples, each	2017-03-14 ~ 2017-03-21	Figs. 11 and 15
sodb10	INC2048	300 samples	2017-03-13 ~ 2017-03-20	Figs. 12(a) and 16(a)
sodb12	INC4096	300 samples	2017-03-02 ~ 2017-03-17	Figs. 12(b) and 16(b)
sodb10	INC8192	261 samples	2017-04-27 ~ 2017-05-21	Figs. 12(c) and 16(c)
sodb12	INC16384	130 samples	2017-04-27 ~ 2017-05-21	Figs. 12(d) and 16(d)

Table 2: Notes on experiment runs used for histograms

INC8192/INC16384 unfortunately stopped in the middle of their runs due to a frozen vnc problem. So couldn't finish 300 samples.

4.1 ET

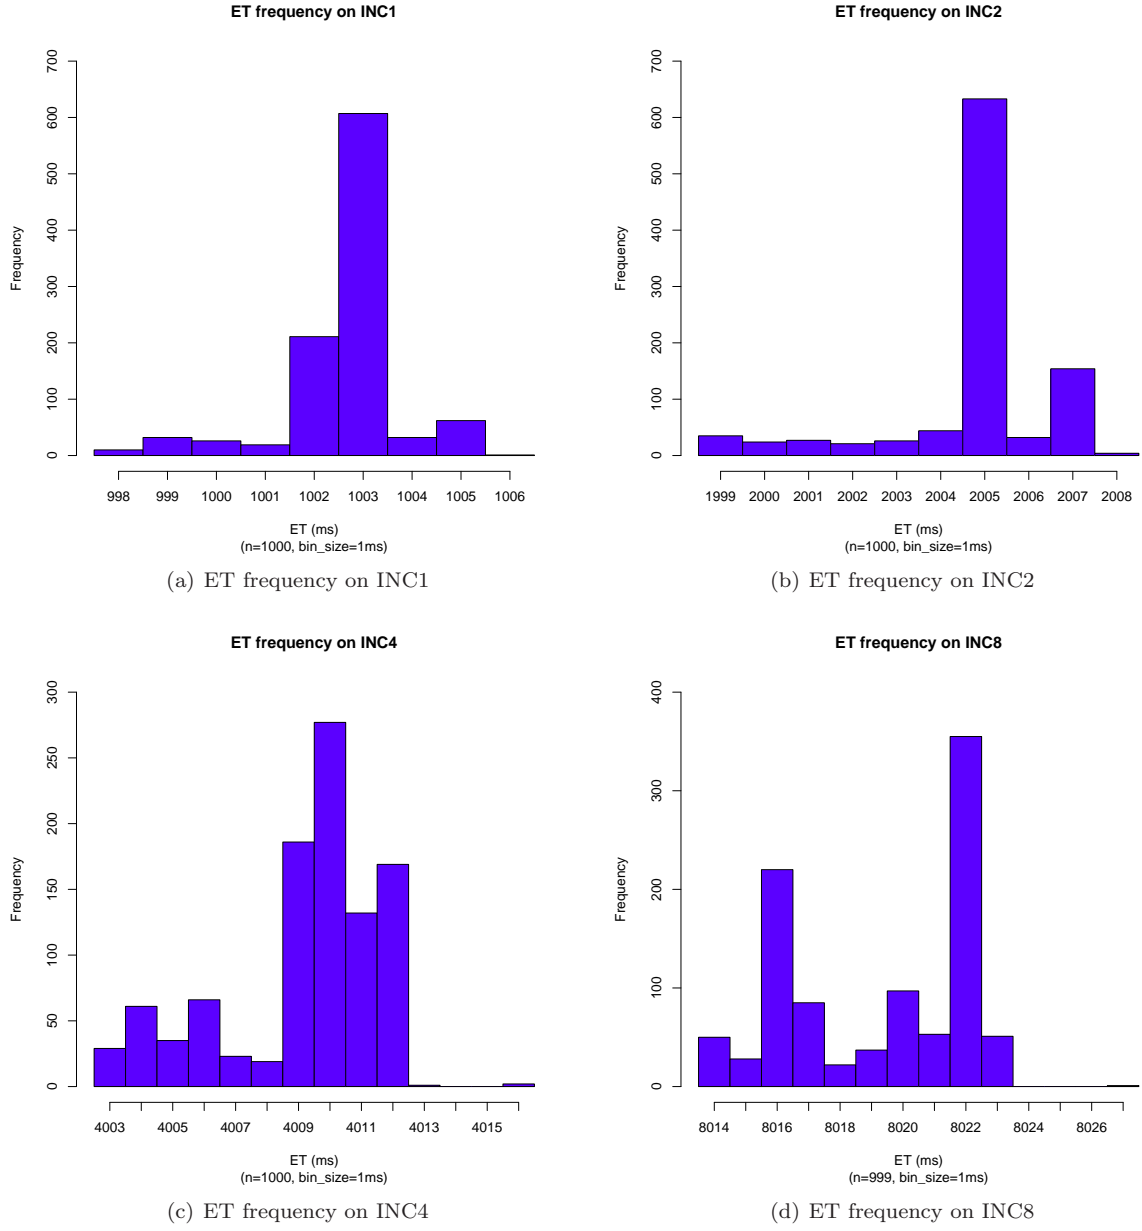





Figure 9: ET Histograms of INC1 ... INC8



(a) ET frequency on INC16




(b) ET frequency on INC32




(c) ET frequency on INC64


Figure 10: ET Histograms of INC16 ... INC64




(a) ET frequency on INC128



(b) ET frequency on INC256




(c) ET frequency on INC512




(d) ET frequency on INC1024


Figure 11: ET Histograms of INC128 ... INC1024




(a) ET frequency on INC2048



(b) ET frequency on INC4096



(c) ET frequency on INC8192



(d) ET frequency on INC16384

Figure 12: ET Histograms of INC2048 ... INC16384

4.2 PT

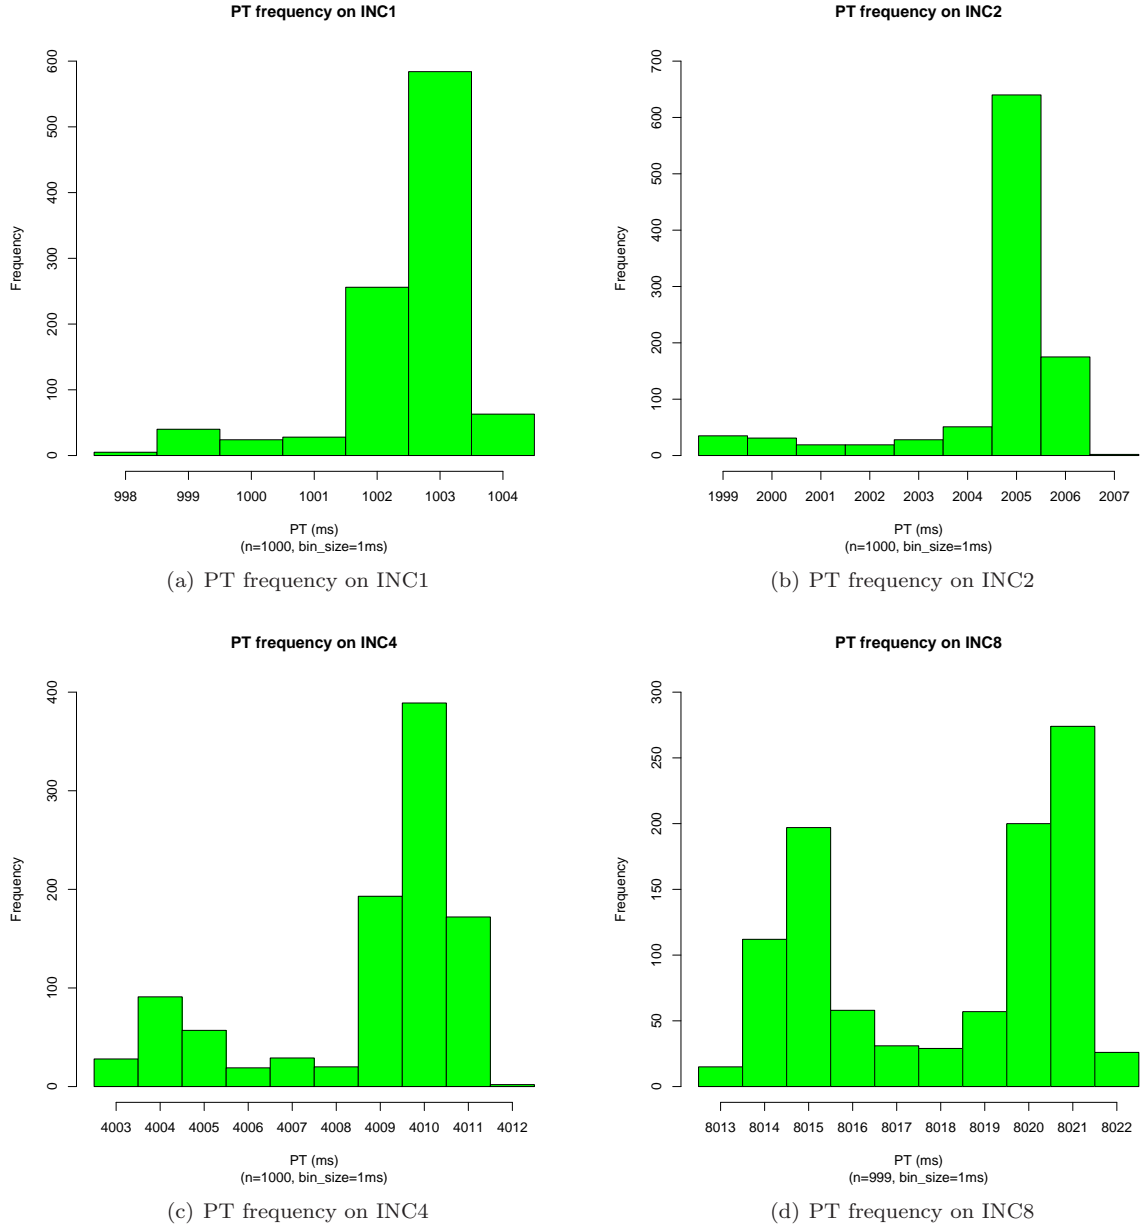





Figure 13: PT Histograms of INC1 ... INC8



(a) PT frequency on INC16




(b) PT frequency on INC32




(c) PT frequency on INC64


Figure 14: PT Histograms of INC16 ... INC64




(a) PT frequency on INC128



(b) PT frequency on INC256




(c) PT frequency on INC512




(d) PT frequency on INC1024


Figure 15: PT Histograms of INC256 ... INC1024




(a) PT frequency on INC2048



(b) PT frequency on INC4096



(c) PT frequency on INC8192



(d) PT frequency on INC16384

Figure 16: PT Histograms of INC2048 ... INC16384

4.2.1 Analysis

In this section we look into what happened inside the peaks observed in a certain histogram. We consider Figure 13(d) for this study. In the figure, we see the peaks at 8015 msec, 8020 msec, and 8021 msec.

Table 3 shows captured daemons and their runtime statistics per bin of figure. Note that bin is at the unit of PT. It appears that the peaks are definitely correlated with (1) appearances of some daemons and (2) times that those daemons co-ran with INC8.

TASK_LEN	BIN (PT)	DAEMON	MIN_PT	MAX_PT	AVG_PT	STD_PT	Counts
INC8	8013	jbd2/md0-8	1	1	1	0	1
INC8	8013	kslowd000	1	1	1	0	1
INC8	8013	md0_raid1	1	1	1	0	17
INC8	8013	proc_monitor	196	200	197.72	1.07	18
INC8	8014	jbd2/md0-8	1	1	1	0	5
INC8	8014	kslowd000	1	1	1	0	35
INC8	8014	kslowd001	1	1	1	0	26
INC8	8014	md0_raid1	1	1	1	0	58
INC8	8014	proc_monitor	196	200	197.31	1.06	95
INC8	8015	java	2	7	4.5	3.54	2
INC8	8015	jbd2/md0-8	1	1	1	0	2
INC8	8015	kslowd000	1	1	1	0	86
INC8	8015	kslowd001	1	1	1	0	89
INC8	8015	md0_raid1	1	1	1	0	18
INC8	8015	proc_monitor	196	200	197.28	1.01	194
INC8	8016	kslowd000	1	1	1	0	36
INC8	8016	kslowd001	1	1	1	0	40
INC8	8016	md0_raid1	1	1	1	0	8
INC8	8016	proc_monitor	196	200	196.45	.95	78
INC8	8017	kslowd000	1	1	1	0	11
INC8	8017	kslowd001	1	1	1	0	10
INC8	8017	md0_raid1	1	1	1	0	3
INC8	8017	proc_monitor	196	200	197.15	1.16	26
INC8	8018	kslowd000	1	1	1	0	13
INC8	8018	kslowd001	1	1	1	0	9
INC8	8018	md0_raid1	1	1	1	0	6
INC8	8018	proc_monitor	196	200	197.24	1.27	29
INC8	8019	jbd2/md0-8	1	1	1	0	3
INC8	8019	kslowd000	1	1	1	0	9
INC8	8019	kslowd001	1	2	1.06	.24	18
INC8	8019	md0_raid1	1	1	1	0	27
INC8	8019	proc_monitor	196	200	197.1	1.18	52
INC8	8020	jbd2/md0-8	1	1	1	0	8
INC8	8020	kslowd000	1	1	1	0	52
INC8	8020	kslowd001	1	1	1	0	57
INC8	8020	md0_raid1	1	1	1	0	91
INC8	8020	proc_monitor	196	200	197.03	1.02	180
INC8	8021	cifs	1	1	1	0	1
INC8	8021	java	2	37	19.5	24.75	2
INC8	8021	kslowd000	1	1	1	0	146
INC8	8021	kslowd001	1	1	1	0	143
INC8	8021	md0_raid1	1	1	1	0	11
INC8	8021	proc_monitor	196	198	197.15	.98	299
INC8	8022	kslowd000	1	1	1	0	20
INC8	8022	kslowd001	1	1	1	0	9
INC8	8022	proc_monitor	196	198	196.07	.37	29

Table 3: Daemons observed from the INC8 run

5 Histograms with 10,000 samples

This section exhibits histograms on two runs of INC, each with 8 and 16 seconds as its task length, having 10,000 repetitions. The detailed description of the base data is from Table 4.

Machine	Task Length (sec)	Description	Experiment Period	Relevant Histograms
sodb9	INC8	10000 samples	2017-03-29 ~ 2017-03-30	Figs. 17(a) and 17(b)
sodb10	INC16	10000 samples	2017-03-29 ~ 2017-03-31	Figs. 17(c) and 17(d)

Table 4: Notes on experiment runs used for histograms




Figure 17: Histograms of INC8 and INC16 with 10,000 samples

6 Further invstigation of the Second Run

This section exhibits user and system time histograms on the second run of INC with its task length increasing from 1 second to 4096 seconds, via SEDONA. The detailed description of the base data is from Table 2.

6.1 User Time

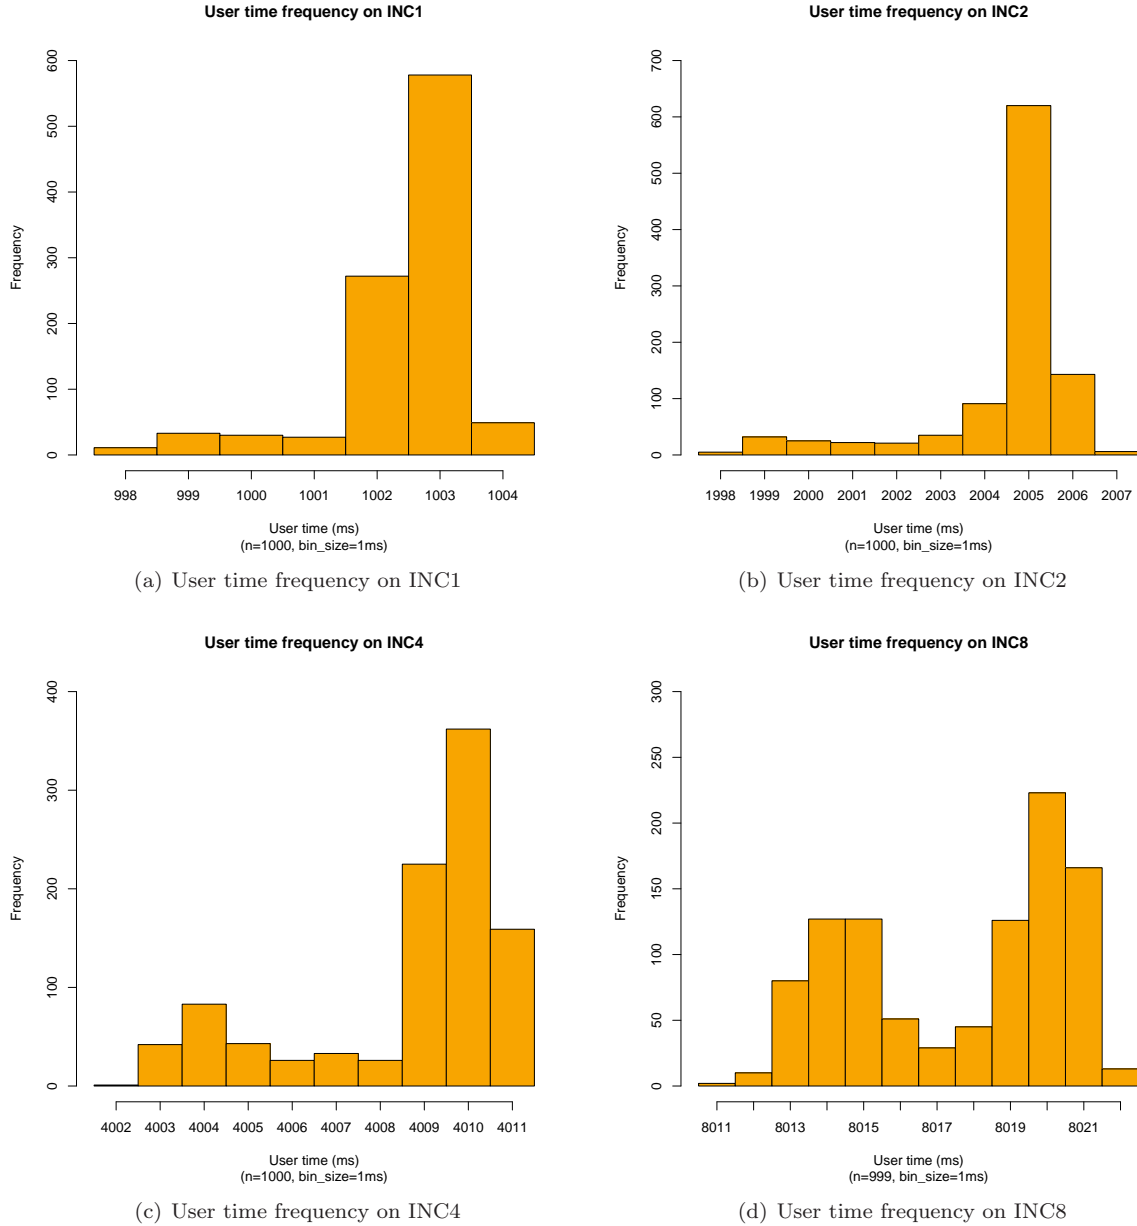





Figure 18: User Time Histograms of INC1 ... INC8



(a) User time frequency on INC16




(b) User time frequency on INC32




(c) User time frequency on INC64


Figure 19: User Time Histograms of INC16 ... INC64




(a) User time frequency on INC96



(b) User time frequency on INC128




(c) User time frequency on INC160




(d) User time frequency on INC192


Figure 20: User Time Histograms of INC96 ... INC192



(a) User time frequency on INC256



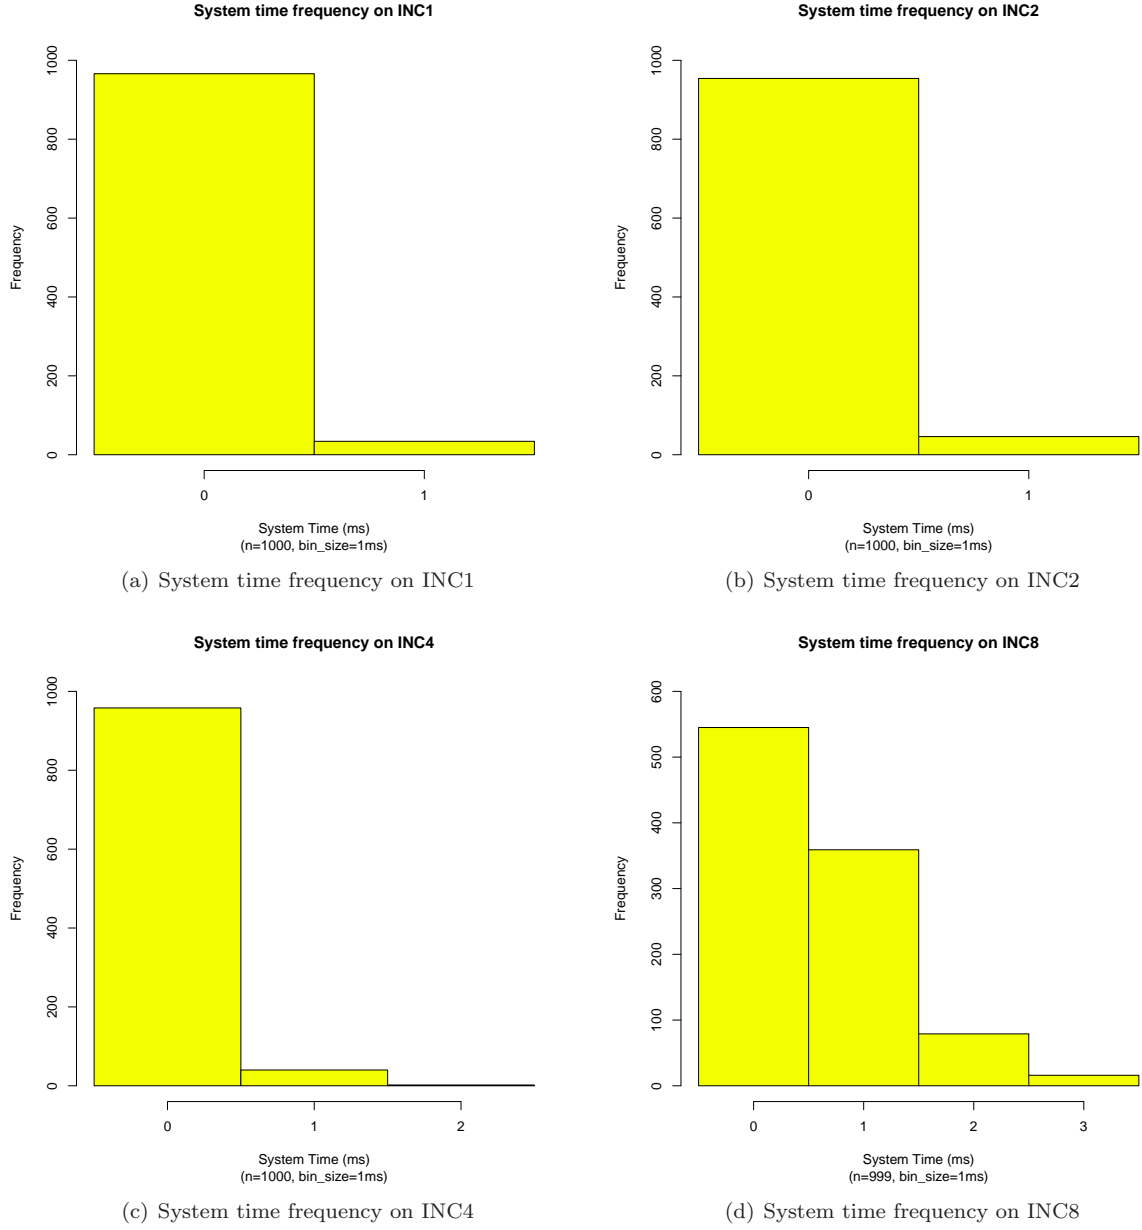
(b) User time frequency on INC512 (with one extreme outlier of 513,259 msec removed)




(c) User time frequency on INC1024


Figure 21: User Time Histograms of INC256 ... INC1024

6.2 System Time






(a) System time frequency on INC16




(b) System time frequency on INC32




(c) System time frequency on INC64


Figure 23: System Time Histograms of INC16 ... INC64




(a) System time frequency on INC128



(b) System time frequency on INC256




(c) System time frequency on INC512




(d) System time frequency on INC1024

Figure 24: System Time Histograms of INC256 ... INC1024



(a) System time frequency on INC2048



(b) System time frequency on INC4096

Figure 25: System Time Histograms of INC2048 and INC4096

6.3 Correlation

	INC2048's u time	INC2048's s time
INC2048's u time	-	0.05
daemon u time	0.56	0.61
daemon s time	0.48	0.57
daemon minor faults	0.53	0.62
daemon major faults	0.55	0.59
daemon read bytes	0.55	0.59
daemon read char	0.56	0.61
daemon read sys calls	0.57	0.63
daemon write bytes	0.57	0.64
daemon write char	0.53	0.62


Table 5: Correlation of user and system time of INC2048 with some daemon measures

	INC4096's u time	INC4096's s time
INC4096's u time	-	-0.30
daemon u time	0.1	0.3
daemon s time	-0.09	0.19
daemon minor faults	0.11	0.32
daemon read char	0.1	0.32
daemon read sys calls	0.11	0.32
daemon write bytes	0	0.26
daemon write char	0.11	0.32


Table 6: Correlation of user and system time of INC4096 with some daemon measures

6.4 Scatter Plots on Some Significant Correlations


The following scatter plots correspond to the correlations bold in Table 6.




(a) User time vs. Sum of all daemon system time



(b) System time vs. Sum of all daemon system time




(c) User time vs. Sum of all daemon minor faults




(d) System time vs. Sum of all daemon minor faults


Figure 26: Scatter plots between measures on INC4096




(a) System time vs. Sum of all daemon read char



(b) System time vs. Sum of all daemon read syscalls



(c) System time vs. Sum of all daemon write bytes



(d) System time vs. Sum of all daemon write char

Figure 27: Scatter plots between measures on INC4096

6.5 Further Investigation of Some Samples

6.5.1 Samples in Fig. 26(c)

proc name (id)	u time	s time	min flt	maj flt	r bytes	r char	r sysc	w bytes	w char	w sysc
INC4096 (3559)	4103920	36	0	0	0	1	1	0	0	0
proc_monitor (25917)	194	6	2	0	0	0	0	0	0	0
md0_raid1 (484)	0	3	0	0	0	0	0	0	0	0
java (3549)	2	1	1093	0	0	11480	20	0	0	0
cifs (1927)	0	1	0	0	0	0	0	0	0	0
kblockd/0 (16)	0	1	0	0	0	0	0	0	0	0
ntpd (28232)	0	1	1	0	0	0	42	4096	7	0

Table 7: Observed values of measures of processes on the leftmost sample in Fig. 26(c)

proc name (id)	u time	s time	min flt	maj flt	r bytes	r char	r sysc	w bytes	w char	w sysc
INC4096 (3559)	4103984	19	0	0	0	1	1	0	0	0
proc_monitor (25917)	190	6	2	0	0	0	0	0	0	0
java (3549)	2	1	1093	0	0	11480	20	0	0	0
md0_raid1 (484)	0	3	0	0	0	0	0	0	0	0
ntpd (28232)	0	0	1	0	0	0	39	4096	7	0
java (4108)	0	0	20	0	0	0	0	0	0	0

Table 8: Observed values of measures of processes on the second rightmost sample in Fig. 26(c)

proc name (id)	u time	s time	min flt	maj flt	r bytes	r char	r sysc	w bytes	w char	w sysc
INC4096 (3559)	4103988	25	0	0	0	1	1	0	0	0
proc_monitor (25917)	194	2	2	0	0	0	0	0	0	0
sshd (3609)	8	4	1382	0	0	512357	400	0	20881	0
bash (3611)	4	1	835	0	0	283911	155	0	136	0
md0_raid1 (484)	0	3	0	0	0	0	0	0	0	0
java (3549)	2	1	1093	0	0	11480	20	0	0	0
cifsd (1927)	0	2	0	0	0	0	0	0	0	0
java (3606)	0	1	22	0	0	0	0	0	0	0
jbd2/md0-8 (497)	0	1	0	0	0	0	0	450560	0	0
kblockd/0 (16)	0	1	0	0	0	0	0	0	0	0
grep (3617)	1	0	311	0	0	5417	11	0	0	0
bash (3612)	0	0	158	0	0	0	0	0	0	0
consoletype (3613)	0	0	127	0	0	1956	6	0	7	0
bash (3614)	0	0	174	0	0	0	0	0	0	0
uname (3615)	0	0	189	0	0	1956	6	0	7	0
sshd (3610)	0	0	425	0	0	22656	29	0	4630	0
sshd (2105)	0	0	14	0	0	0	0	0	594	0
bash (3618)	0	0	170	0	0	0	0	0	0	0
id (3619)	0	0	225	0	0	4352	12	0	2	0
ntpd (28232)	0	0	1	0	0	0	42	4096	7	0
bash (3616)	0	0	131	0	0	0	0	0	61	0

Table 9: Observed values of measures of processes on the rightmost sample in Fig. 26(c)

6.5.2 Samples in Fig. 27(c)

proc name (id)	u time	s time	min flt	maj flt	r bytes	r char	r sysc	w bytes	w char	w sysc
INC4096 (3559)	4103981	14	0	0	0	1	1	0	0	0
proc_monitor (25917)	194	2	2	0	0	0	0	0	0	0
md0_raid1 (484)	0	5	0	0	0	0	0	0	0	0
java (3549)	2	1	1093	0	0	11480	20	0	0	0
flush-9:0 (3548)	0	2	0	0	0	0	0	0	0	0
ntpd (28232)	0	0	1	0	0	0	42	4096	7	0
java (4585)	0	0	20	0	0	0	0	0	0	0

Table 10: Observed values of measures of processes on the leftmost sample in Fig. 27(c)

proc name (id)	u time	s time	min flt	maj flt	r bytes	r char	r sysc	w bytes	w char	w sysc
INC4096 (3559)	4103958	44	0	0	0	1	1	0	0	0
proc_monitor (25917)	194	2	2	0	0	0	0	0	0	0
sshd (4877)	10	4	1382	0	0	519568	392	0	20868	0
sshd (4886)	10	3	1383	0	0	515456	393	0	20868	0
grep (4888)	6	1	994	0	0	287034	154	0	136	0
grep (4879)	3	2	991	0	0	286801	159	0	136	0
java (3549)	2	1	1093	0	0	11480	20	0	0	0
md0 RAID1 (484)	0	3	0	0	0	0	0	0	0	0
sshd (4878)	2	1	424	0	0	22403	26	0	4268	0
jbd2/md0-8 (497)	0	2	0	0	0	0	0	696320	0	0
flush-9:0 (3548)	0	2	0	0	0	0	0	0	0	0
cifs (1927)	0	1	0	0	0	0	0	0	0	0
kernel (16)	0	1	0	0	0	0	0	0	0	0
grep (4892)	0	1	309	0	0	5417	11	0	0	0
grep (4883)	1	0	310	0	0	5417	11	0	0	0
id (4894)	0	0	226	0	0	4352	12	0	2	0
bash (4884)	0	0	164	0	0	0	0	0	0	0
id (4885)	0	0	226	0	0	0	12	0	2	0
sshd (2105)	0	0	28	0	0	0	0	0	1188	0
sshd (4887)	0	0	425	0	0	22403	26	0	4268	0
ntpd (20232)	0	0	1	0	0	0	42	4096	7	0
bash (4880)	0	0	165	0	0	0	0	0	0	0
uname (4890)	0	0	190	0	0	1956	6	0	7	0
bash (4891)	0	0	128	0	0	0	0	0	61	0
java (4874)	0	0	20	0	0	0	0	0	0	0
bash (4893)	0	0	164	0	0	0	0	0	0	0
uname (4881)	0	0	191	0	0	1956	6	0	7	0
bash (4889)	0	0	165	0	0	0	0	0	0	0

Table 11: Observed values of measures of processes on the second rightmost sample in Fig. 27(c)

proc name (id)	u time	s time	min flt	maj flt	r bytes	r char	r sysc	w bytes	w char	w sysc
INC4096 (3559)	4103948	45	0	0	0	1	1	0	0	0
proc_monitor (25917)	190	10	2	0	0	0	0	0	0	0
md0 RAID1 (484)	0	5	0	0	0	0	0	0	0	0
jbd2/md0-8 (497)	0	3	0	0	0	0	0	471040	0	0
java (3549)	2	1	1093	0	0	11480	20	0	0	0
cifs (1927)	0	2	0	0	0	0	0	0	0	0
flush-9:0 (3548)	0	1	0	0	0	0	0	0	0	0
java (5022)	1	0	21	0	0	0	0	0	0	0
ntpd (28232)	0	0	1	0	0	0	42	4096	7	0

Table 12: Observed values of measures of processes on the rightmost sample in Fig. 27(c)


7 Histograms on Consecutive INC1024 Runs

This section exhibits histograms on (three) consecutive runs of INC1024 via SEDONA. The detailed description of the base data is from Table 2.


Machine	Task Length (sec)	Description	Experiment Period	Relevant Histograms
sodb9	INC1024	300 samples, each	2017-04-12 ~ 2017-04-23	Figs. 28 and 29

Table 13: Notes on experiment runs used for histograms


7.1 ET



(a) ET frequency on INC1024-Run1




(b) ET frequency on INC1024-Run2




(c) ET frequency on INC1024-Run3

Figure 28: ET Histograms of Three Consecutive INC1024 Runs


7.2 PT



(a) PT frequency on INC1024-Run1



(b) PT frequency on INC1024-Run2



(c) PT frequency on INC1024-Run3

Figure 29: PT Histograms of Three Consecutive INC1024 Runs

References

- [1] Sabah Currim, Richard T. Snodgrass, Young-Kyoon Suh, and Rui Zhang, “DBMS Metrology: Measuring Query Time”, in *ACM Transactions on Database Systems*, 42(1):3:1–42(+8), 2017.